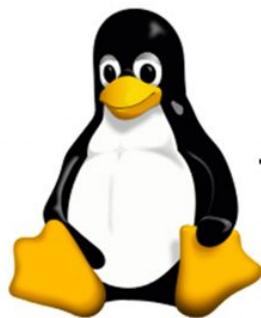


# Transfert de style & Generative Adversarial Networks

## Tout un art au service de l'art !



+



=



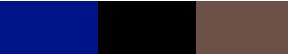
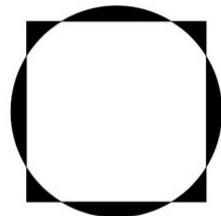


# TDS NoBlaBla : Style Transfer and Generative Adversarial Networks

Julien Guillaumin

[julien.guillaumin@imt-atlantique.net](mailto:julien.guillaumin@imt-atlantique.net)

IMT Atlantique (ex Télécom Bretagne)

- 
- 
- Student at **IMT Atlantique**, Brest.
    - Engineering degree - Computer Vision
  - Deep Learning Engineer at **Lighton.io**
    - Optical co-processor for Machine Learning
  - **SummerIA.fr**, Deep Learning Tutor
  - **Continental**, Deep Learning for ADAS (intern)
  - **Thales Services**, Large-scale Image Processing with Spark (intern)
  - Talks and MeetUps
    - **PyCon France 2016**
    - **Toulouse Data Science, Paris AI, Data Science Brest, Machine Learning Aix-Marseille**
    - **Airbus Defence and Space, Quantmetry, Ercom**
- 

# Outline

- **Introduction to Deep Learning**
  - Basics of Machine Learning
  - Machine Learning vs Deep Learning
  - Convolutional Neural Networks - CNNs
- **Style Transfer with Neural Networks**
  - Perceptual loss : content loss + style loss
  - Optimization-based method: first steps in neural style transfer
  - Train CNNs to perform neural style transfer
- **Generative Adversarial Networks - GANs**
  - How to generate realistic images ?
  - Two networks, two players: Hi Game Theory !
  - GANs for semi-supervised learning and domain adaptation
  - Can we understand the latent space ?
  - Applications: Super-resolution



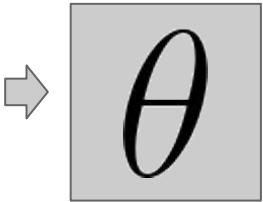
# Introduction to Deep Learning

# Basics of Machine Learning

- **Field of Artificial Intelligence**
  - **Develop learnable algorithms**
    - Learn from data
    - To resolve complex tasks
  - **Many tasks :**
    - Natural Language Processing
    - Image classification
    - Object segmentation
- Two major phases :**
- **Training**
    - From training data
    - Adjust internal parameters
    - Goal : find generalization !
  - **Inference**
    - New data (not seen)
    - Evaluation / Production

# Basics of Machine Learning

## Case of Image Classification



$$f_{\theta}(x^{(i)})$$

$$y^{(i)} \rightarrow \text{bird}$$

$$E = -\log (P(Y = y^{(i)} | x^{(i)}, \theta))$$

*negative log-likelihood  
E : error for this example*

$$\mathbb{P}(Y = \text{truck} | x^{(i)}, \theta)$$

$$\mathbb{P}(Y = \text{bird} | x^{(i)}, \theta)$$

0

0

0

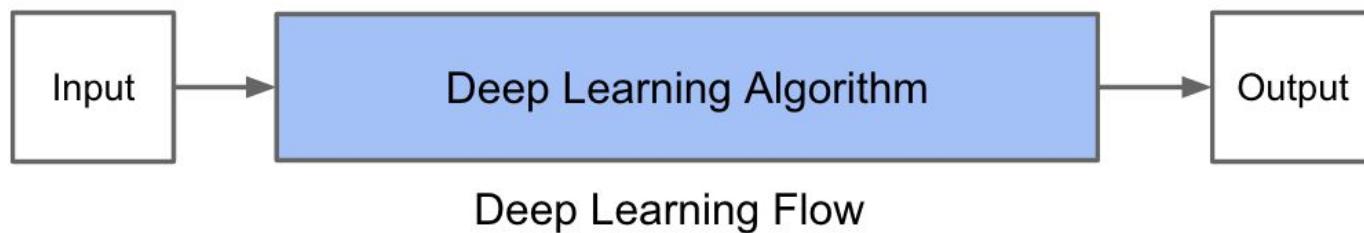
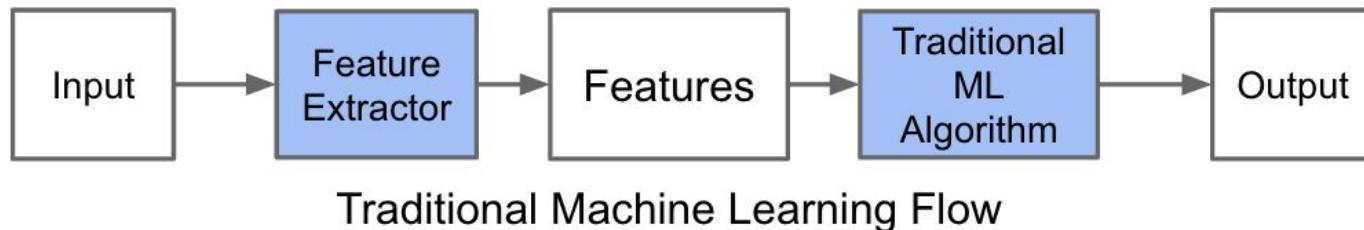
$$\mathbb{P}(Y = \text{plane} | x^{(i)}, \theta)$$

0.12

0.81

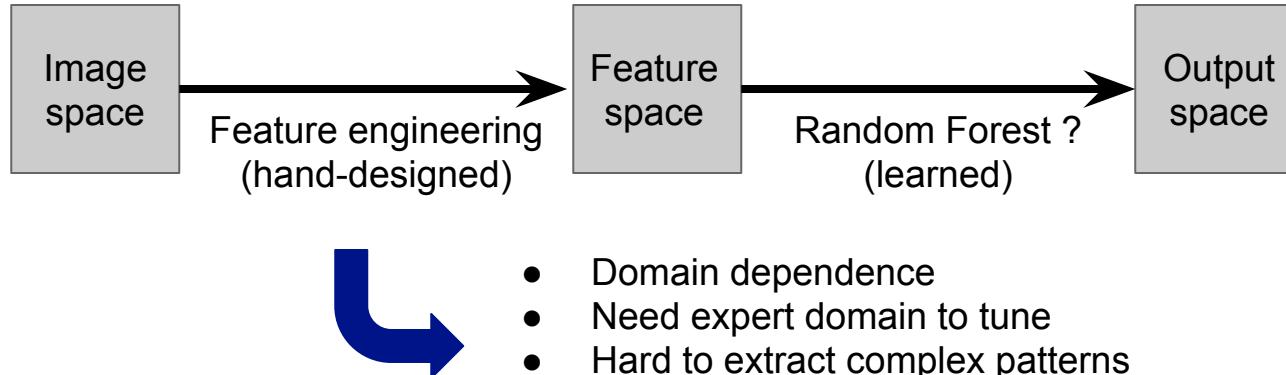
0.05

# Machine Learning vs. Deep Learning



# Machine Learning vs Deep Learning

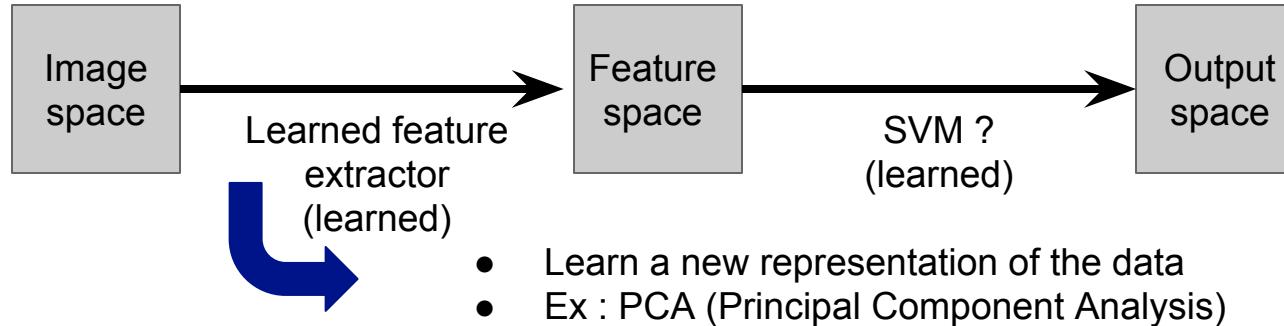
## Machine Learning approach



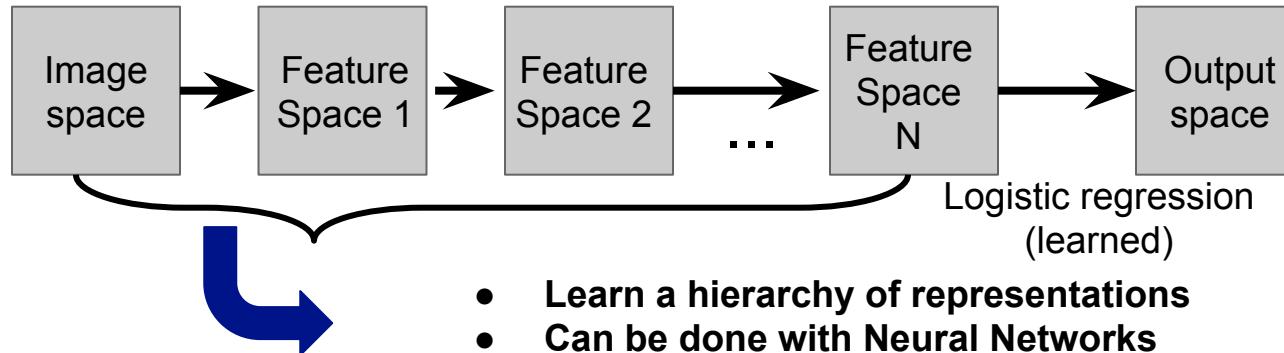
**For images :** HOG features, SIFT methods, Histograms, LBP features, ....

# Machine Learning vs. Deep Learning

## Representation Learning approach

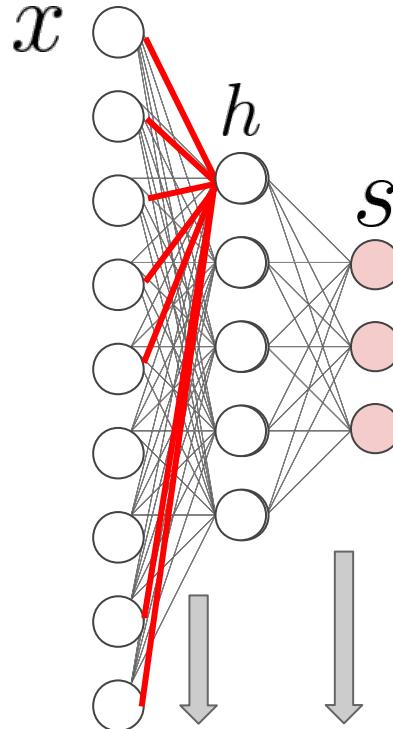


## Deep Representation Learning approach



# Deep Neural Networks - DNNs

- Biologically inspired
- Representation as vectors
- Learn to perform vector transformations
- Weighted sum + biases
- Activation function



$$\theta = \{W_1, b_1, W_2, b_2\}$$

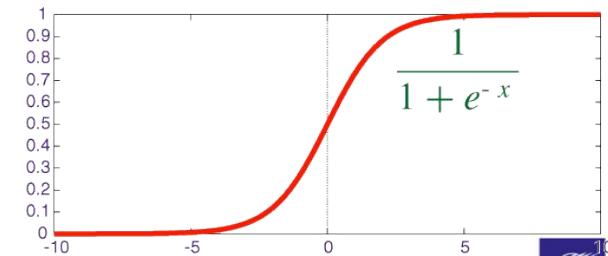
1. Weighted sum + biases
2. Activation function

Weights and biases are learned !

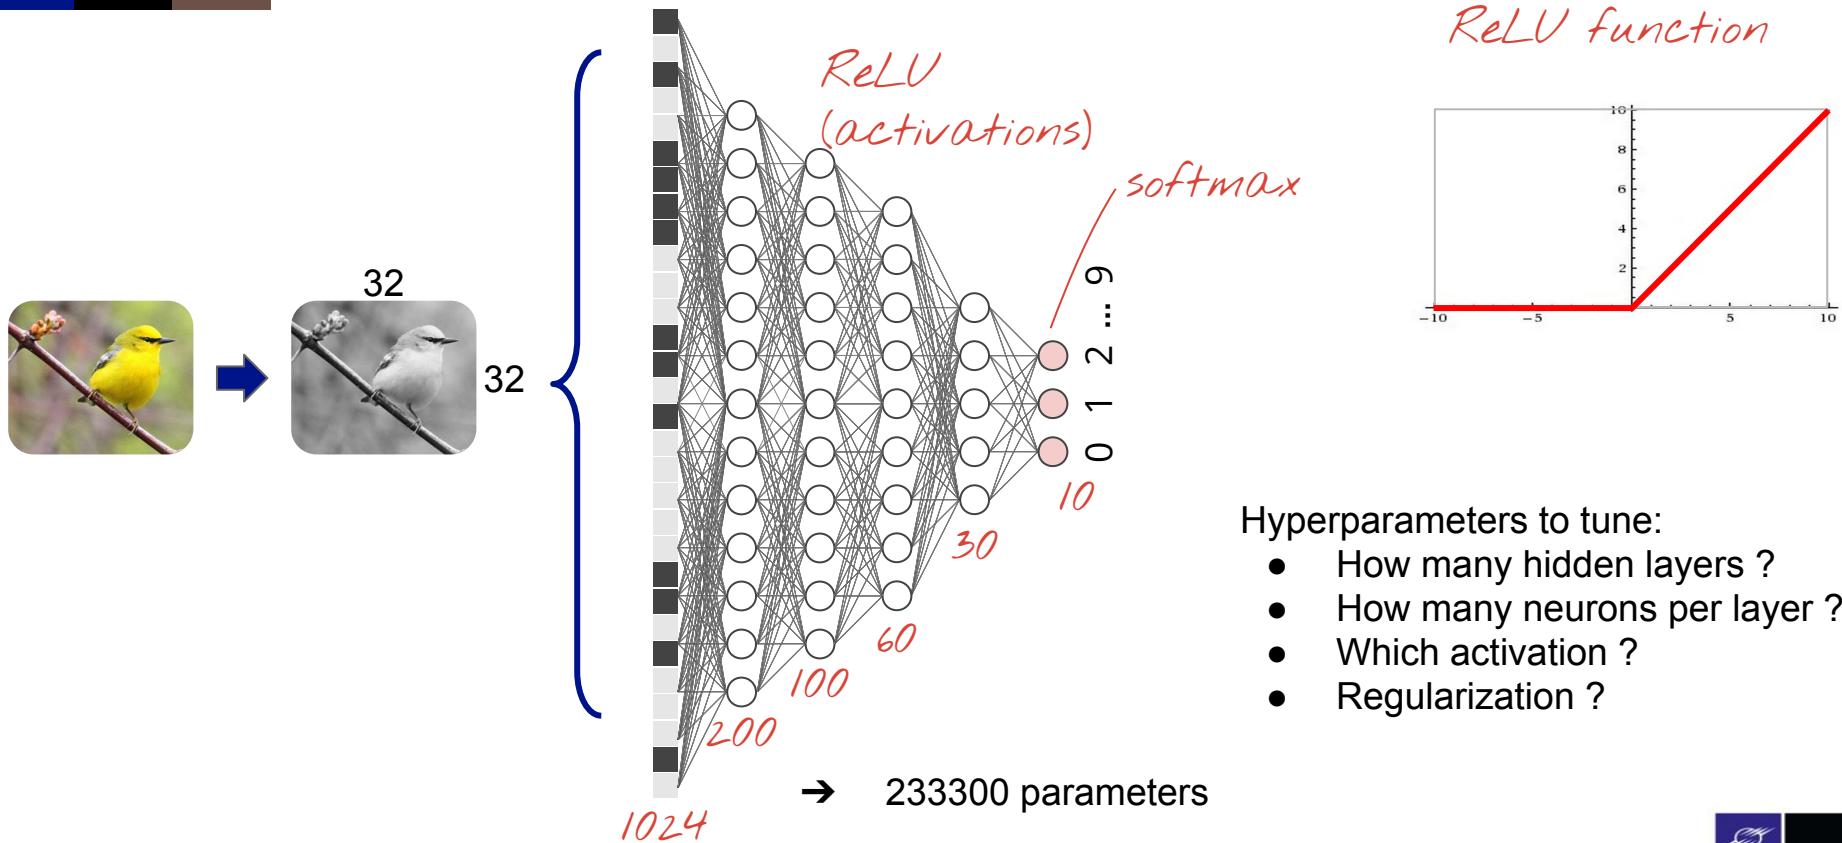
$$h = act(W_1x + b_1)$$
$$s = softmax(W_2x + b_2)$$

Activation function :

- Sigmoid
- Tanh
- ReLU !



# Deep Neural Networks - DNNs



Hyperparameters to tune:

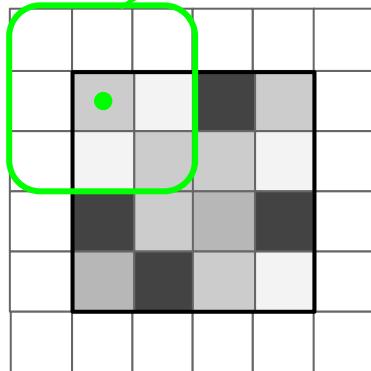
- How many hidden layers ?
- How many neurons per layer ?
- Which activation ?
- Regularization ?

# Convolutional Neural Networks - CNNs

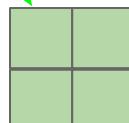
Intuition for CNNs :

- Keep 2D representation
- High correlation between adjacent pixels
- Weight sharing
- Many hyper-parameters :
  - kernel size, padding, stride, with bias ?

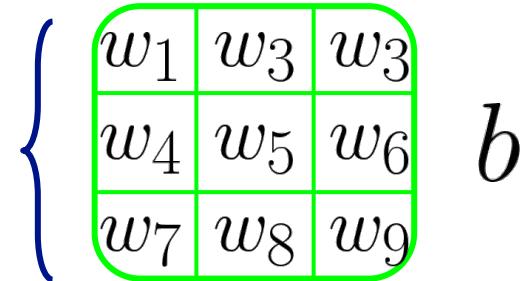
$x : [4,4]$  + zero padding



- kernel 3x3
- padding = 'same'
- stride = 2



To learn



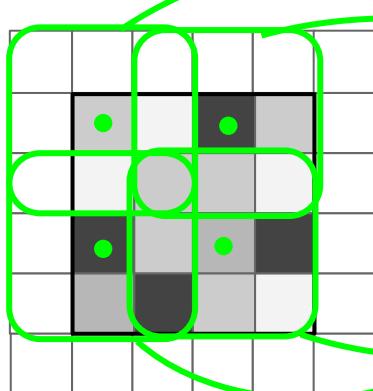
# Convolutional Neural Networks - CNNs

Intuition for CNNs :

- Keep 2D representation
- High correlation between adjacent pixels
- Weight sharing

- Many hyper-parameters :
  - kernel size, padding, stride, with bias ?

$x : [4,4] + \text{zero padding}$



- kernel 3x3  
- padding = 'same'  
- stride = 2

To learn

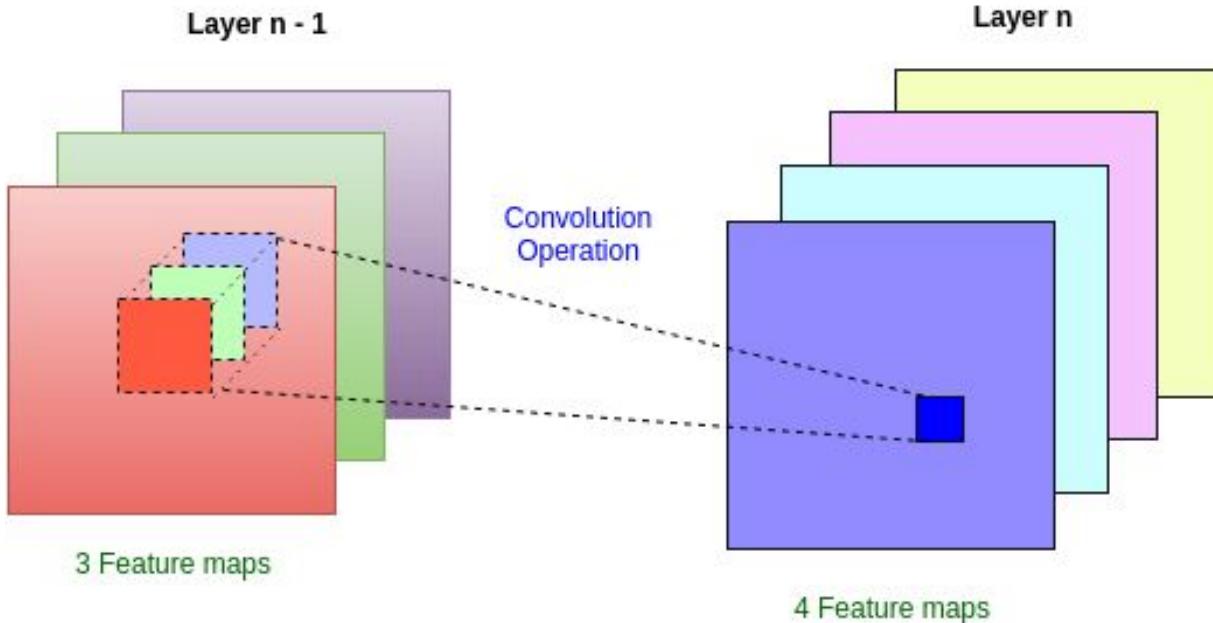
$w_1$	$w_3$	$w_3$
$w_4$	$w_5$	$w_6$
$w_7$	$w_8$	$w_9$

$b$

$$h = \text{act}(x * W + b)$$

# Convolutional Neural Networks - CNNs

New representation is composed of “Feature Maps”

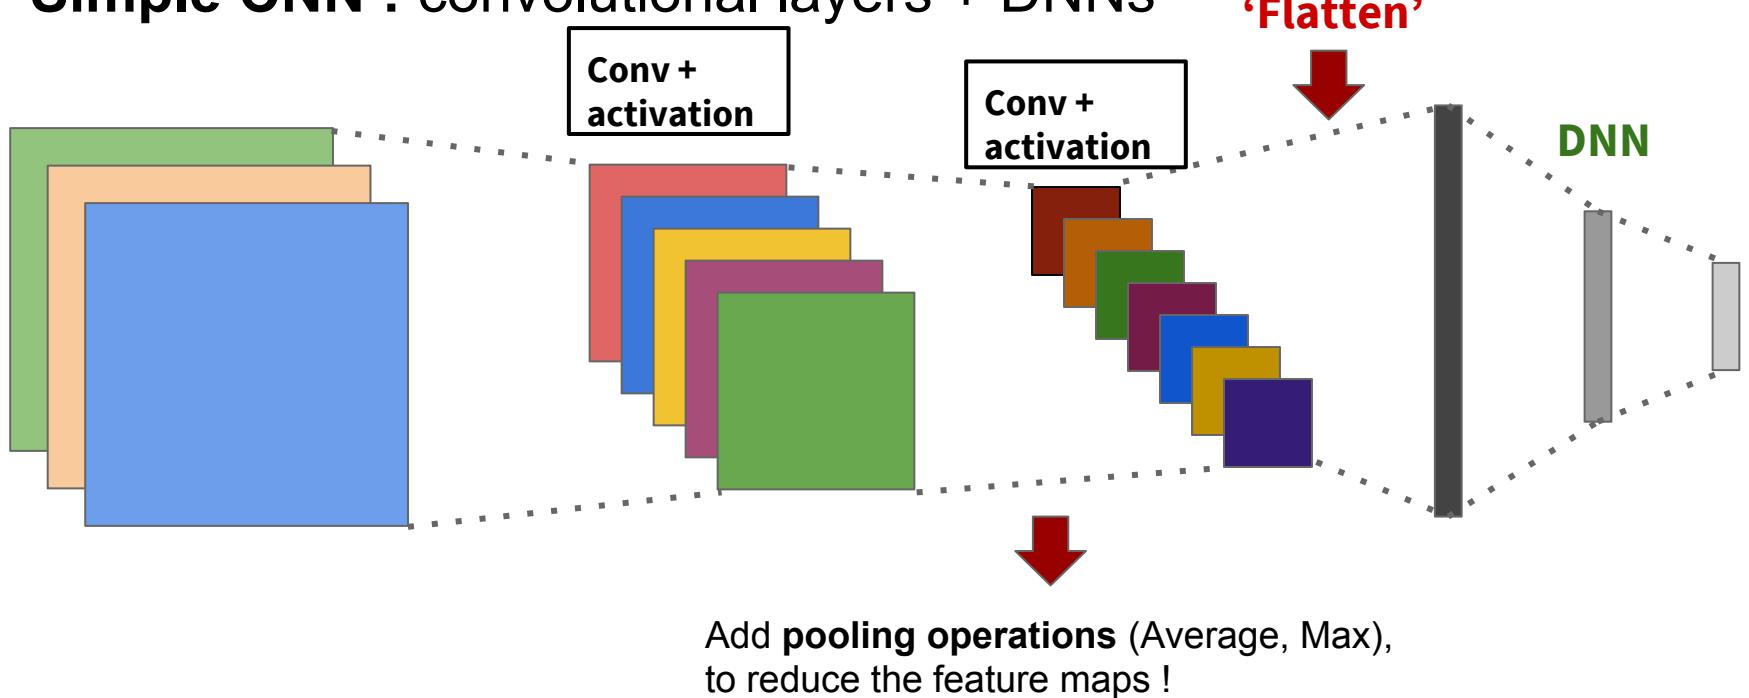


Here :

- 4 kernels to create 4 feature maps
- from 3 feature maps (RGB images, for example)
- $(3 \times 3 \times 3 + 1) \times 4 : 112$  parameters !

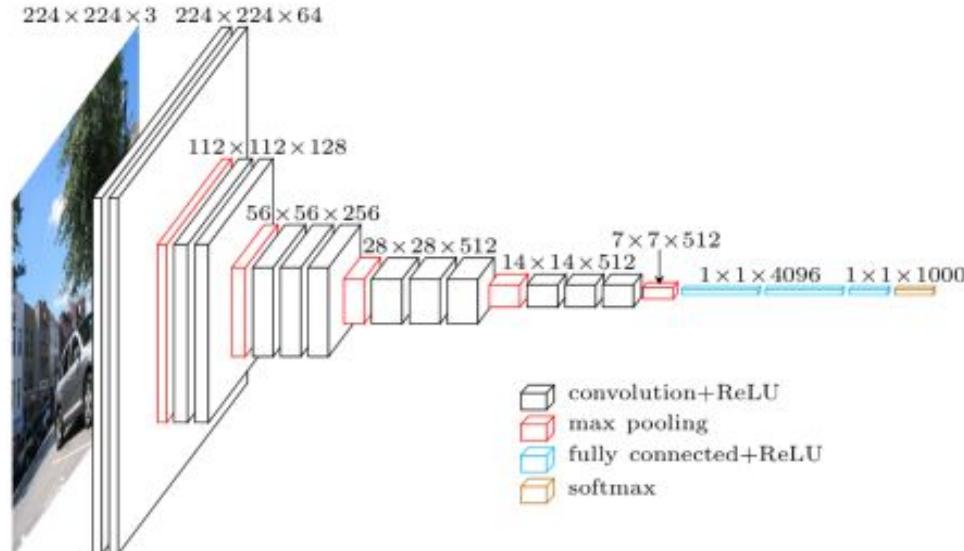
# Convolutional Neural Networks - CNNs

Simple CNN : convolutional layers + DNNs



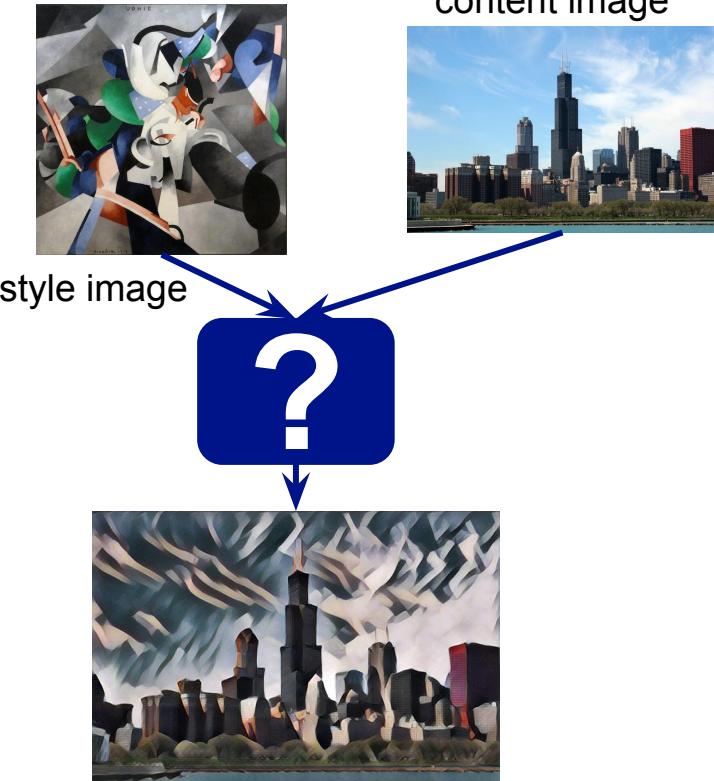
# Deep Network : VGG-16 [1]

- Simple : no **inception modules[2]** or **residual connections[3]**
- Trained for **image classification** on ImageNet[4] (1000 classes)
- State of the art in 2014 (92.7% top-5 test accuracy)
- 138,357,544 parameters (10% conv weights, 90% FC layers)



# Neural Style Transfer

# Neural Style Transfer: Motivations



- **Generative task**
  - From an image, generate a new one
- **Introduction to more complex tasks**
  - Super-resolution and colorisation
- **CNNs understanding is required**
  - Hierarchy of representations
  - Feature spaces

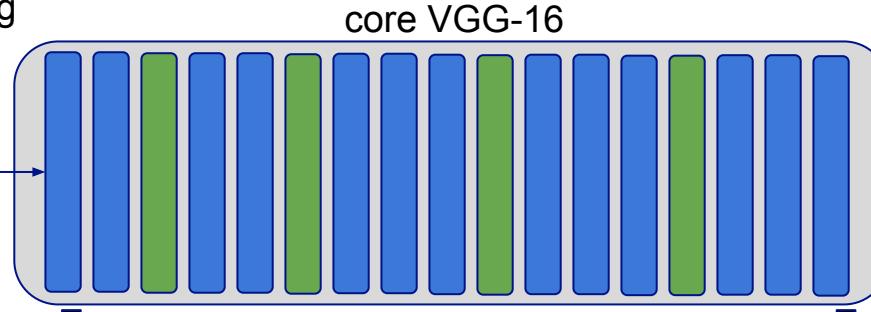
stylized image with content

# CNN visualization

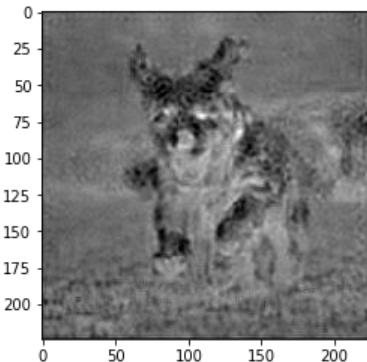
- Convolution + ReLU
- Pooling



preprocessing



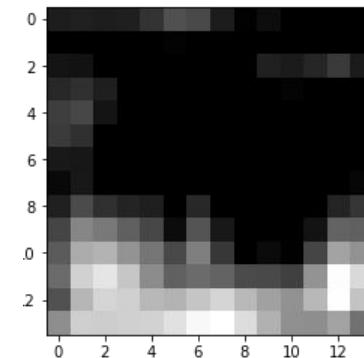
conv1\_1 (224x224x64)



From low-level to  
high-level feature spaces

Additional visualization methods :

- Deep Dream approach [5]
  - Optimization-based
- Zeiler & Fergus [6]
  - Transposed convolutions and unpooling operations

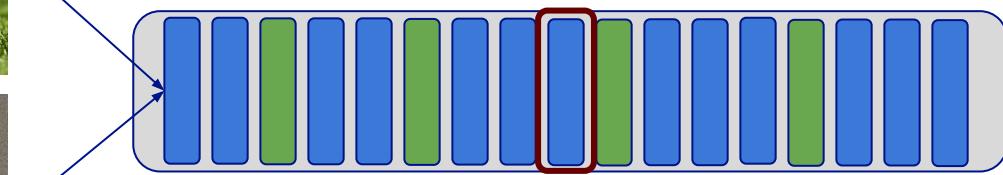


# Content Representation/Reconstruction



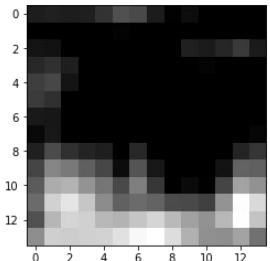
$\Phi_j(x)$  : activations of the  $j$ th layer

conv3\_3 : 56x56x256



Eg :

$\Phi_j(x)[40]$



Fixed VGG-16

$$\hat{x} = \operatorname{argmin}_x \|\Phi_j(x) - \Phi_j(x_c)\|_2^2$$

$x$

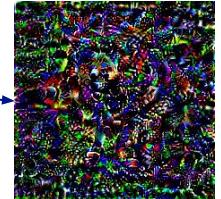
- Goal : find an image with the same activations at a given layer (all feature maps)
- Optimization problem, start from a random image

# Content Representation/Reconstruction

- gradient descent optimization on input image, network does not change
- loss = MSE on feature maps, 1000 iterations, Adam ( $\text{lr}=2.0$ )
- low-level : input image is correctly reconstructed, with pixel-level details
- high-level : only content is preserved

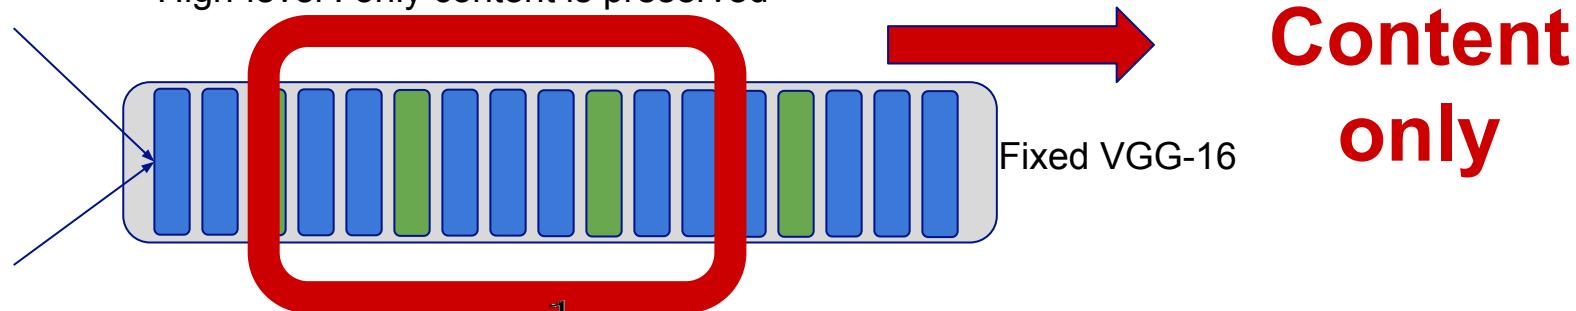
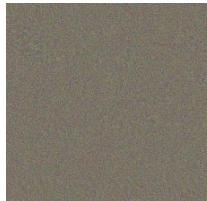
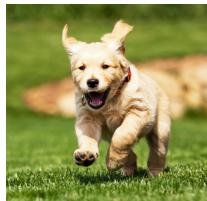


Fixed VGG-16



# Content Representation/Reconstruction

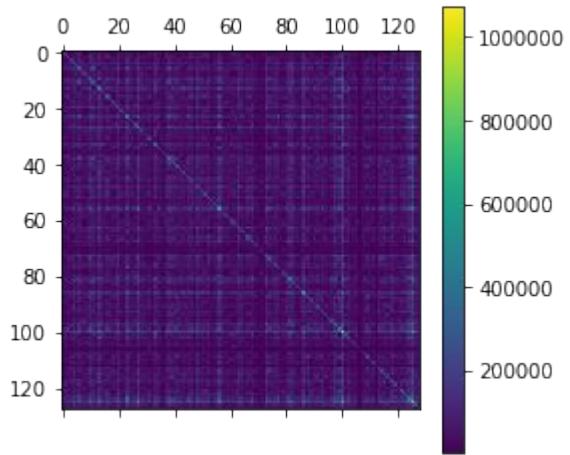
- From a random image, reconstruct the feature maps obtained with a normal image, on a specific layer
- Gradient descent optimization on image input, network does not change
- Loss = MSE on feature maps, 1000 iterations, Adam ( $\text{lr}=2.0$ )
- Low-level : input image is correctly reconstructed, with pixel-level details
- High-level : only content is preserved



$$L_c(x, x_c) = \frac{1}{C_j H_j W_j} \|\phi_j(x_c) - \phi_j(x)\|_2^2$$

# Style Representation/Reconstruction

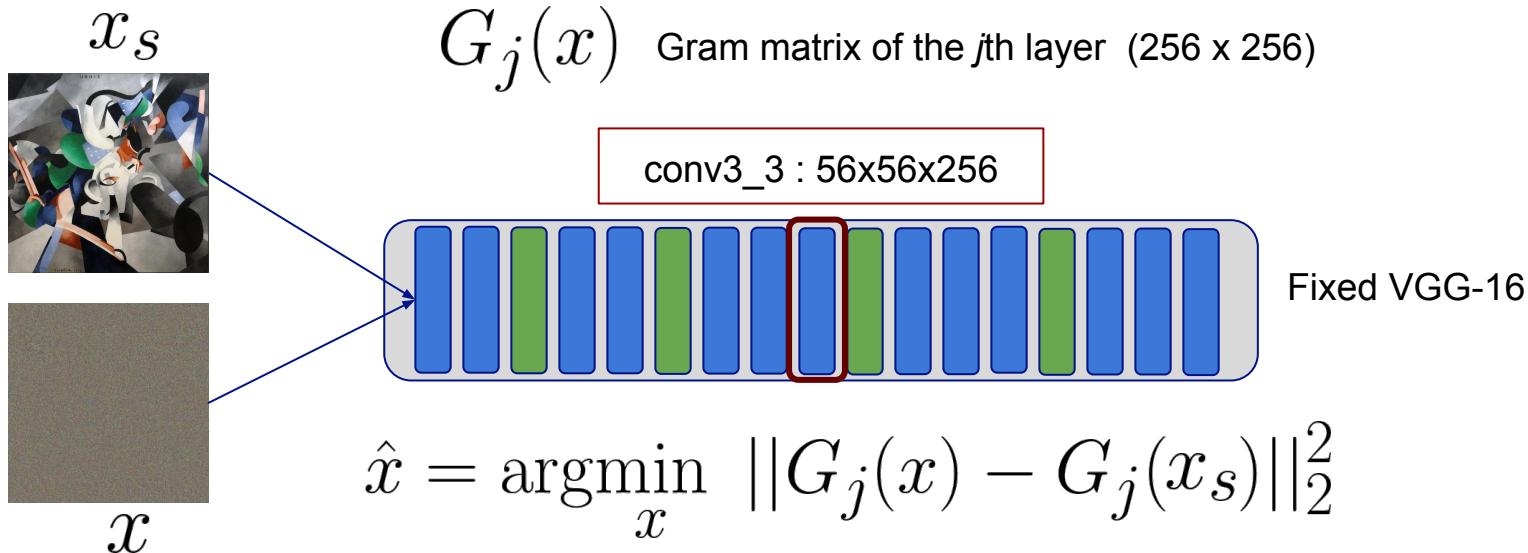
- Needs more complex statistics on feature maps : **Gram matrix**
  - Second-order statistics
  - Can capture texture information, no spatial information
- For a given layer  $j$  with  $C_j$  feature maps of size  $(W_j, H_j)$
- The Gram matrix is a  $(C_j, C_j)$  matrix :



$$G_j(x)_{c_1, c_2} = \mathbb{E}[ \Phi_j(x)[c_1] * \Phi_j(x)[c_2] ]$$

- Where  $*$  is an element-wise operation between 2 feature maps (Hadamard product)
- Contains the correlation between every pair of feature maps

# Style Representation/Reconstruction



- Goal : To find an image with the same Gram matrix for a given layer
- Optimization problem: Start from a random image

# Style Representation/Reconstruction

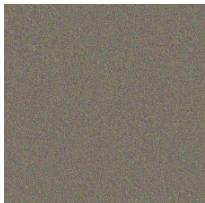
$x_s$



- Gradient descent optimization on image input, network is freezed
- Loss = MSE on feature maps, 1000 iterations, Adam (lr=2.0)
- Low-level : Small and simple patterns
- High-level : More complex patterns

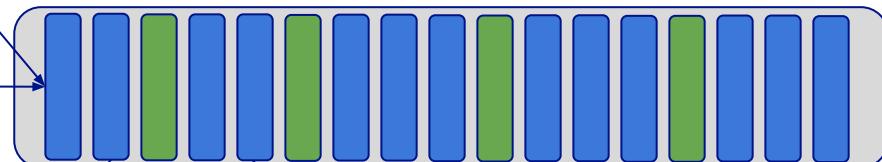
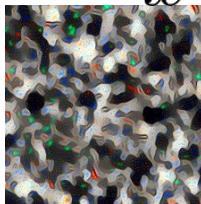
$$L_s(x, x_s) = \sum_j \frac{\lambda_j}{C_j^2} \|G_j(x_s) - G_j(x)\|_2^2$$

Fixed VGG-16



$x$

$\hat{x}$



$\hat{x}$



$\hat{x}$



$\hat{x}$



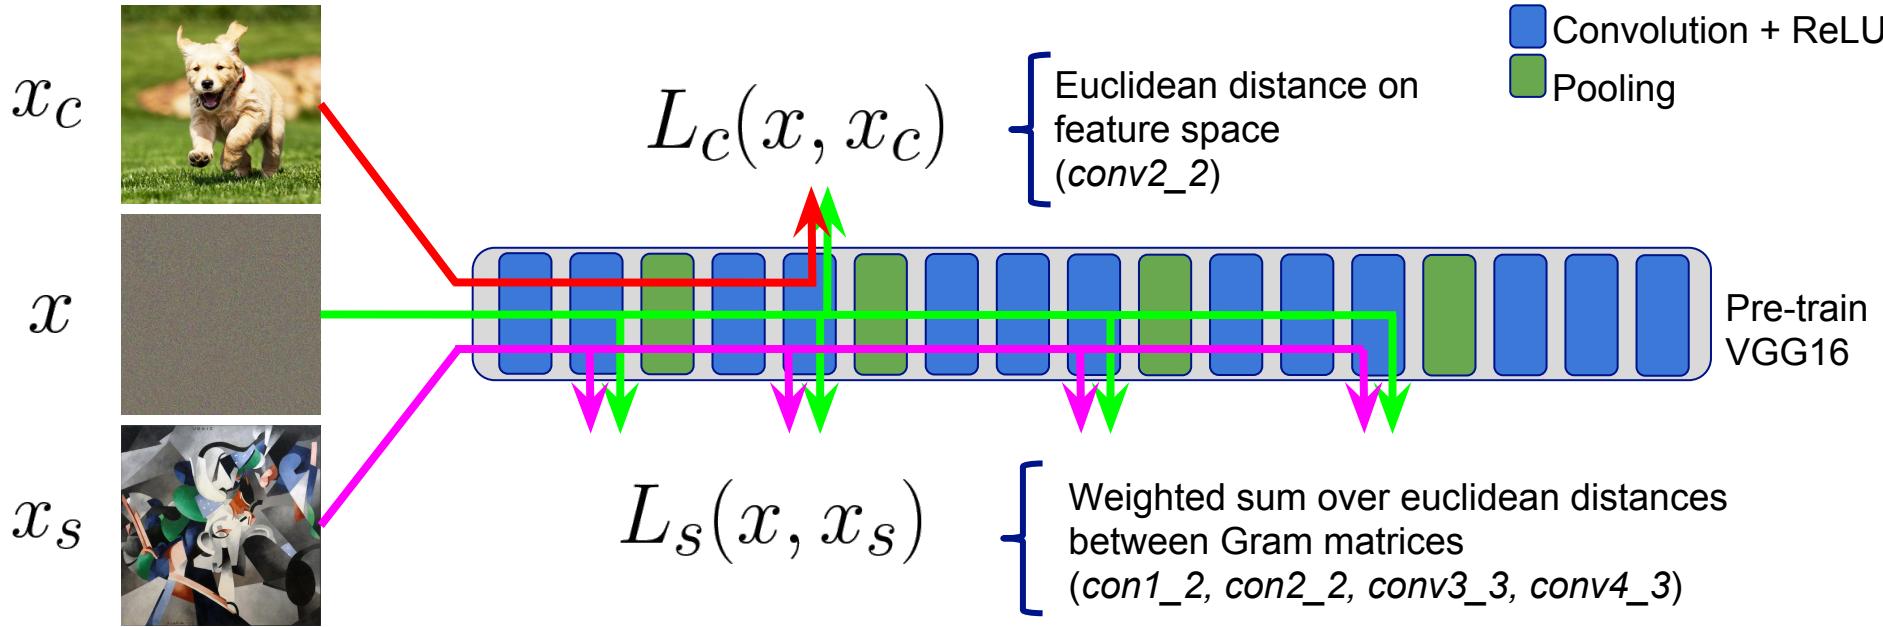
$\hat{x}$



# Content & Style Representations

- Content is preserved in high level features
- Style is present in second-order statistics in low and medium levels
- **Content and Style are separable**
- ***content\_loss* and a *style\_loss* are defined**
- Combine style and content from different images is possible, via feature extraction learned within a VGG network, trained on a generic image classification task !

# Mix content & style via specific losses



$$L(x, x_s, x_c) = \alpha L_s(x, x_s) + \beta L_c(x, x_c)$$

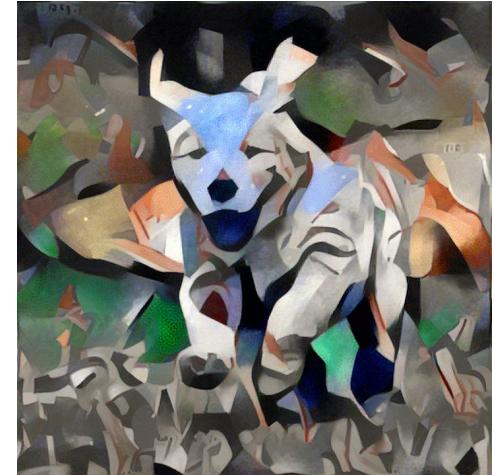
Perceptual loss and method defined in [7]



# Optimization process

- Compute *content\_target* (feature maps) with *content\_image*
- Compute *style\_target* (Gram matrices) with *style\_image*
- Start from a random image (*input\_image*)
- Optimization process :
  - Compute *content\_loss* and *style\_loss* with targets + *input\_image*
  - Minimize this loss by modifying *input\_image*
  - Possible thanks to gradient-descent method (like Adam)

# TensorFlow implementation



$it = [100, 400, 800, 3000]$

# Results

- Produce high-quality images
- Easy to tune effects (more content ? more style ?)
- Any input/output size
- Running time (1000 #iter)
  - GPU (GTX 1070) : ~ 5 min (1920 CUDA cores)
  - CPU (i7-7700K) : ~ 150 min (4 cores x 2 threads)
- Avoid any real-time applications
- But perceptual loss (content+style) is defined

# Improvements

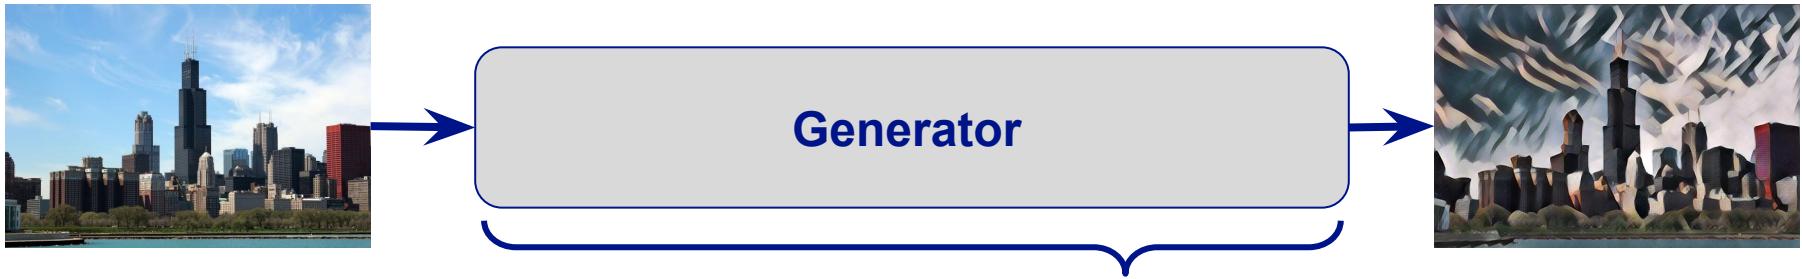
- Time dependency for video transformation (see [8])
- Change optimizer : L-BFGS !
- Tune weights between style and content loss
- Start from : content image? style image? noisy image? or a mix?
- Color constraint : preserve color from content image ! (see [9])



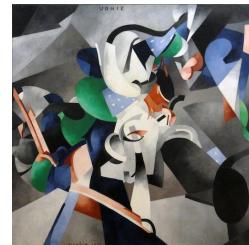
from : [github.com/tensorflow/magenta](https://github.com/tensorflow/magenta)

# Feed-forward method [10, 11]

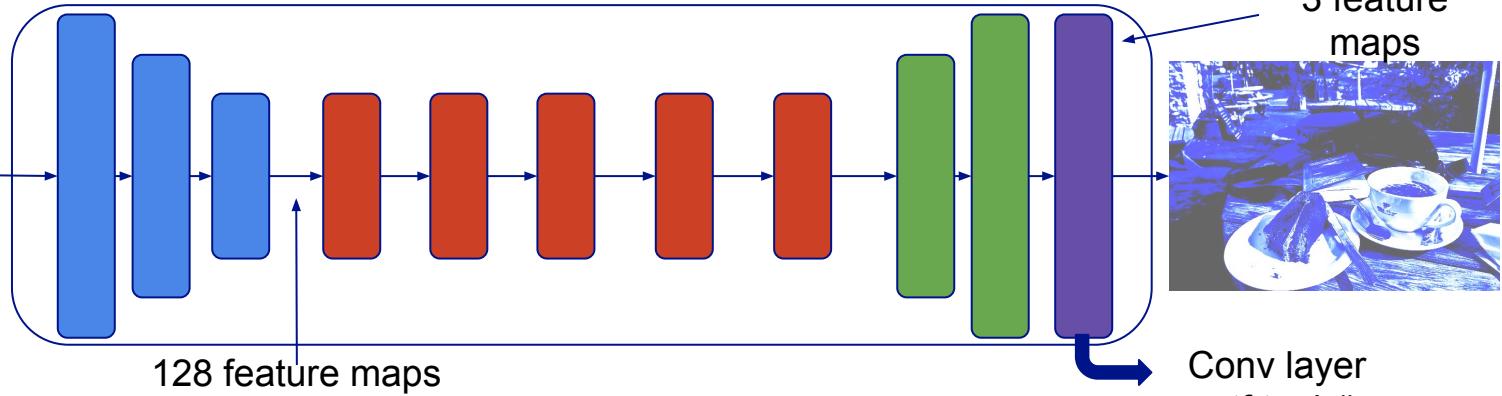
- Train a network to obtain a stylized image in one pass as an output
- Used for one specific style (fixed)



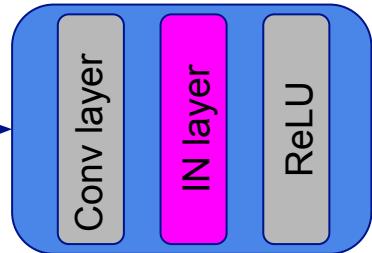
- Trained to add this style
- With a dataset of content images
- Same input/output size



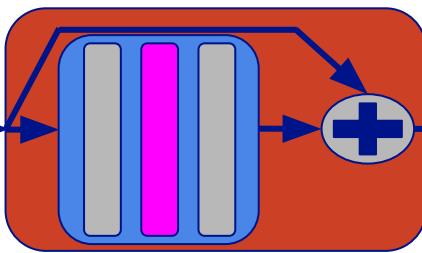
# Architecture of the generator ?



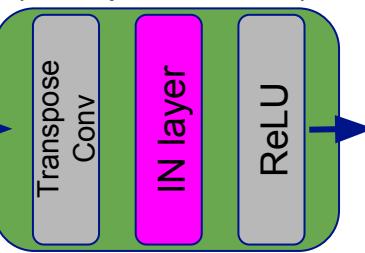
Conv\_block



Residual\_block



Deconv\_block  
(transposed conv)



**Instance Normalization [11]**  
Variant of Batch Normalization

# How to train a generator ?



$$L_c(x_c, \theta)$$



$$L_s(x_s, \theta)$$



- Train with batch of content images
- Minimize total loss w.r.t. theta

# Need a dataset of content images

- COCO dataset[12], about 80k images
- Only 1 style image

## Training process (loop) :

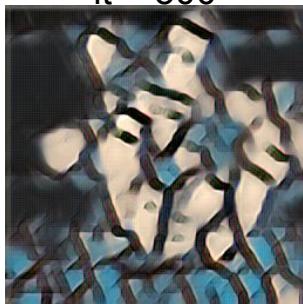
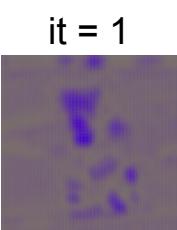
- Take a batch of samples from COCO
- Pass this batch through the generator to get generated images
- Compute *style\_loss* between the generated images and the style image
- Compute *content\_loss* between the generated images and the original ones
- Minimize the *total\_loss* by updating the weights from the generator

## Training information :

- Adam optimizer ( $\beta_1=0.05$ )
- Only 20k iterations (with *batch\_size*=4)
- For 512x512x3:
  - **Training time (on GTX 1070) : 10 hours**
  - **Inference time : 330 ms (GTX 1070)**

# Results and improvements

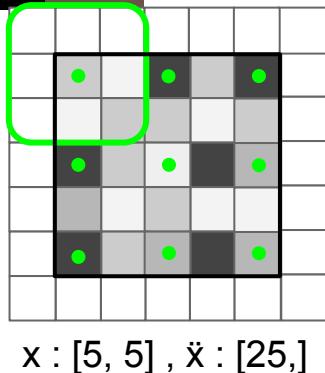
- With a new content image :



- Learn to apply only one style (with fixed style/content levels!!)
- In [13] (ICLR 2017) :
  - Add '**Conditional Instance Normalization**'
  - Learn to apply a fixed set of styles (until 64)
  - Can learn quickly a new style (incremental learning)
- Use **resized convolutions**[14] instead of transposed convolutions : Improves quality
- Add **variational loss** to encourage spatial smoothness
- Now : Universal/Arbitrary Style Transfer ! [15, 16]



# Conv vs. Transposed Conv



$y : [3, 3], \hat{y} : [9,]$

An output tensor  $y$  of size 3x3 with values: 
 

.	.	.
.	.	.
.	.	.

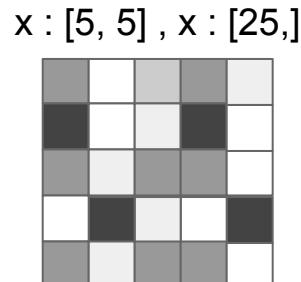
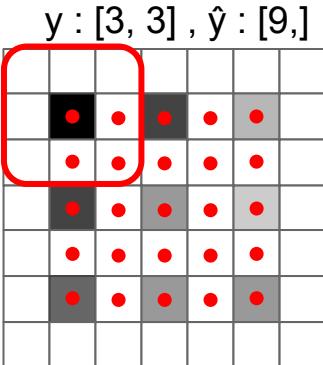
Conv2d, kernel size : 3x3, stride=2, padding='same'

$$w * x = y$$

$$M \ddot{x} = \ddot{y}$$

$M : [9, 25]$

$w_1$	$w_3$	$w_3$
$w_4$	$w_5$	$w_6$
$w_7$	$w_8$	$w_9$



TransposeConv2d, kernel size 3x3, stride=2, padding="same "

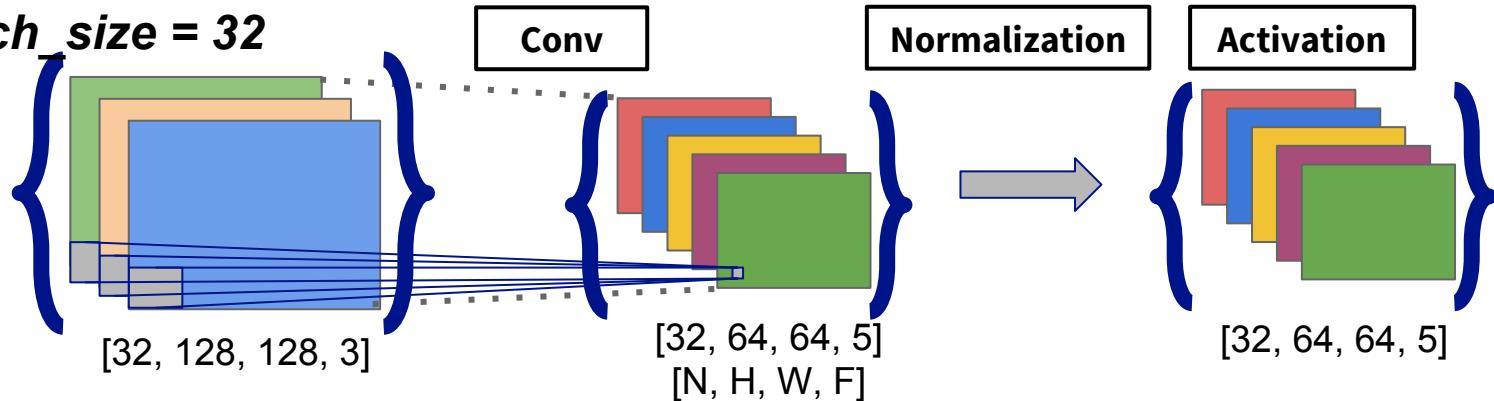
Conv2d, kernel size 3x3, stride=1, "internal zero padding", padding="valid",

$$M^\top \hat{y} = \hat{x}$$

More info about resized conv and transposed conv : <https://distill.pub/2016/deconv-checkerboard/>

# BatchNorm vs. InstanceNorm

batch size = 32



$$x[..., f] \leftarrow \alpha_f \frac{x[..., f] - \mathbb{E}[x[..., f]]}{\sigma(x[..., f])} + \beta_f$$

**BatchNorm :**  
channel-wise  
Discriminative tasks !

$$x[n, ..., f] \leftarrow \alpha_f \frac{x[n, ..., f] - \mathbb{E}[x[n, ..., f]]}{\sigma(x[n, ..., f])} + \beta_f$$

**InstanceNorm :**  
(sample,channel)-wise  
Generative tasks !

# Conditional Instance Normalization : Add meta-data to your CNNs !

Add conditions on  $\{\alpha, \beta\}$  within an Instance Normalization layer :

- Traffic Sign classification :  $\{\alpha, \beta\}_{night}$  vs  $\{\alpha, \beta\}_{day}$
- SAR images :  $\{\alpha, \beta\}(inc\_angle)$



[13] : Conditional Instance Normalization applied to Style Transfer

- **64 styles** with 1 generator and 64 sets of normalization parameters
- **direct interpolation** with the learned normalization parameters to create new styles

# Some results

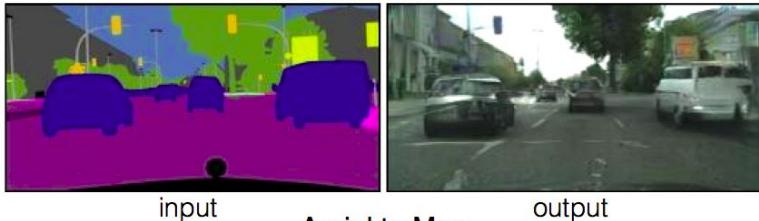




# Generative Adversarial Networks- GANs

# Some “generative” tasks

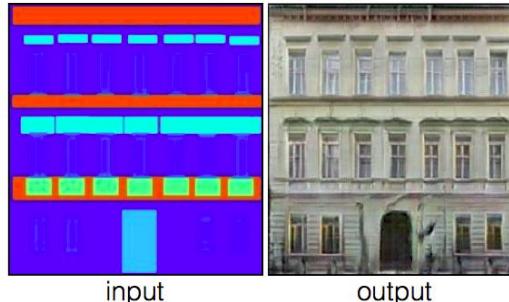
Labels to Street Scene



input

output

Labels to Facade



input

output

BW to Color



input

output

Aerial to Map



input

output

Day to Night



input

output

Edges to Photo



input

output

Source :<https://phillipi.github.io/pix2pix/>

Paper [17]

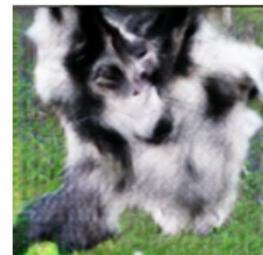
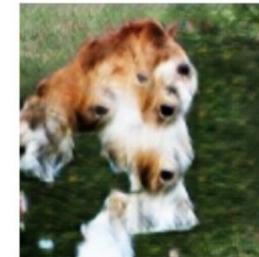
# How to generate realistic images ?

**Task:** given a dataset, generate samples following a distribution similar to the dataset  $p_{data}$

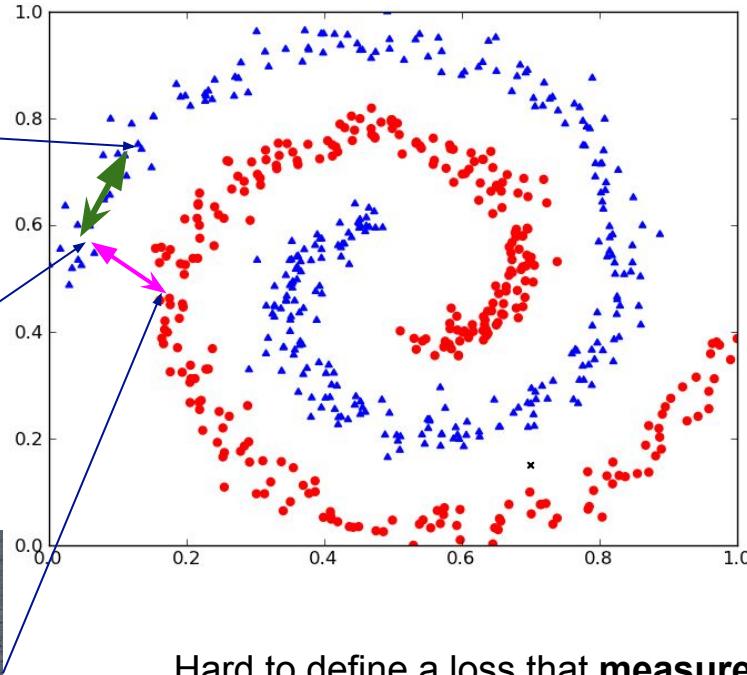
Which loss to use ?

- **MSE** (Mean Squared Error) on image space
- **Total Variation Loss** (impose smoothness)
- **Feature Matching** (MSE on feature maps)
  - Perceptual loss (cf Style Transfer)

→ Blurred images, non-realistic images



# Find the Manifolds of ‘realistic images’ ?



Ships vs Planes manifolds !

Main issue in Machine Learning :

- How to define a good loss for a given task ??

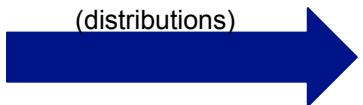
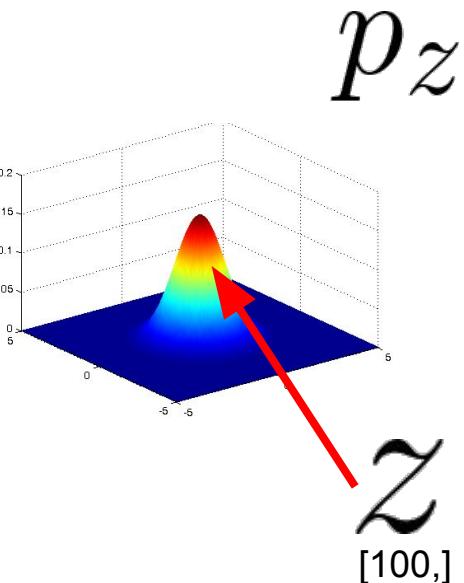
MSE for image generation ?

- Does not capture the concepts
- Distance on low-level representation (pixel-level) !

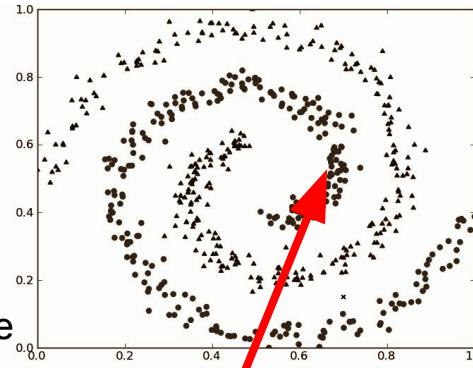
Hard to define a loss that **measure the photorealism** ?

→ **LEARN THIS LOSS WITH NEURAL NETS**

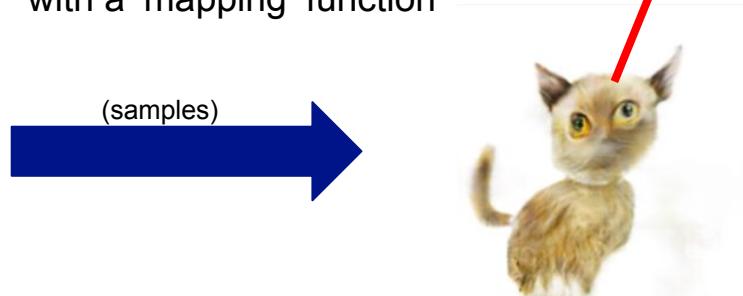
# How to generate cats : Meow generator !



- start from a random noise
- to a realistic image
- in the manifold of cats !
- with a 'mapping' function



$p_{data}$



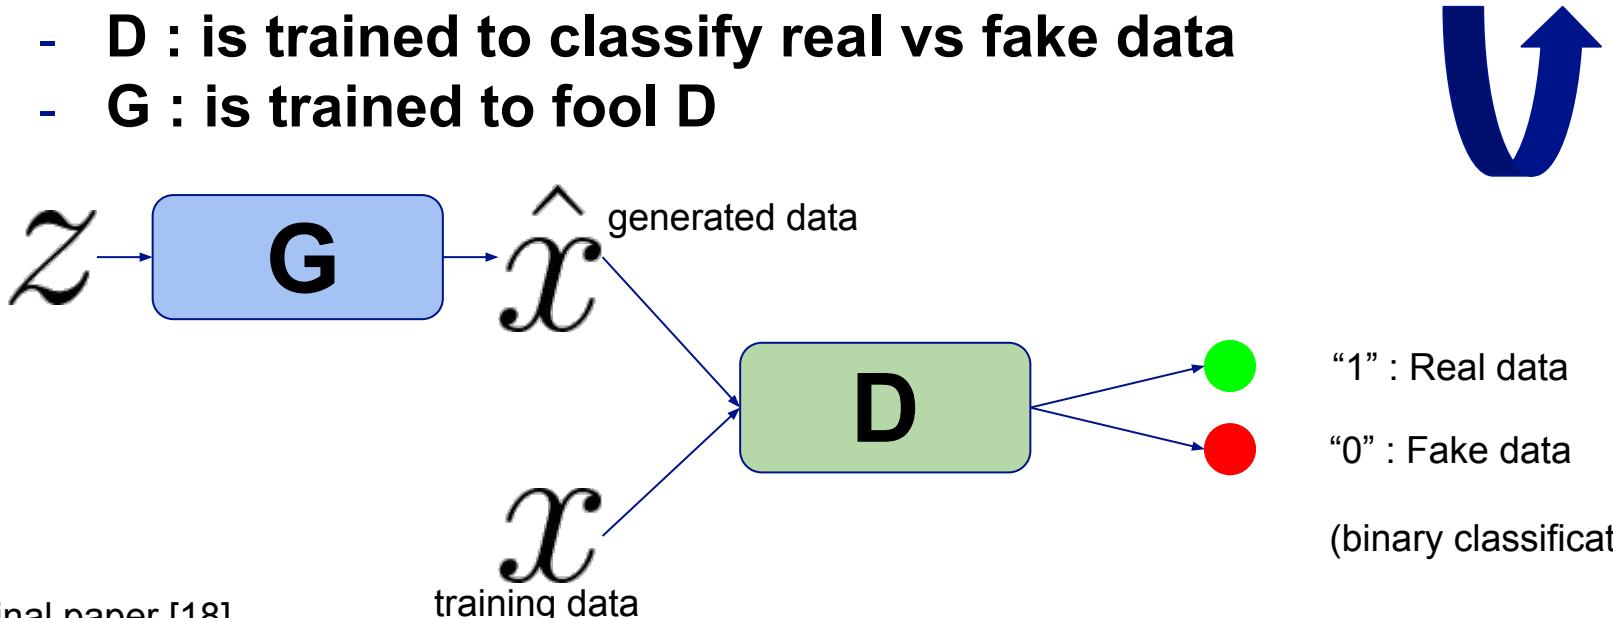
$$\hat{x} = G_\theta(z)$$

[256, 256, 3]

# Generative Adversarial Networks (GANs)

General framework : Generator(G) + Discriminator(D)

- G : generates data from a latent space (noise)
- D : is trained to classify real vs fake data
- G : is trained to fool D



Original paper [18]

# GANs in equations

$$\min_G \max_D \left[ \mathbb{E}_{x \sim p_{data}} [\log(D(x))] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))] \right]$$

min/max game : **Game Theory** - 2 agents : 2 neural networks

- equivalent to minimize the Jensen-Shannon divergence between  $(p_{data} || p_\theta)$
- **Nash equilibrium** :  $D(x) = \frac{1}{2}$  and  $p_\theta = p_{data}$

- Learn an **implicit distribution**, through the generator :

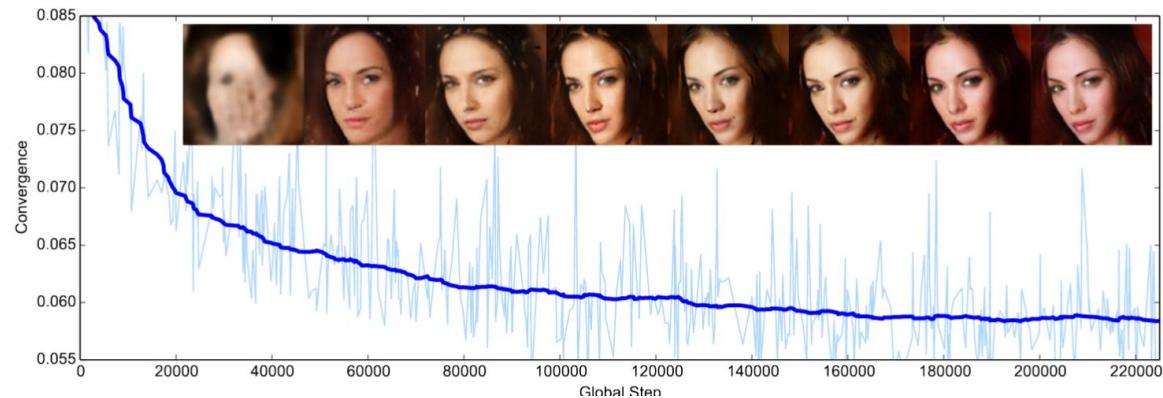
$$z \sim p_z$$
$$x = G_\theta(z) \sim p_\theta$$

# In practice : how to train GANs ?

$$L_D = -\mathbb{E}_{x \sim p_{data}} [\log(D(x))] - \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$$
$$L_G = -\mathbb{E}_{z \sim p_z} [\log(D(G(z)))]$$

## Simultaneously training for D and G:

- train G to fool D with a batch of z
- train D to detect samples from G or from the dataset



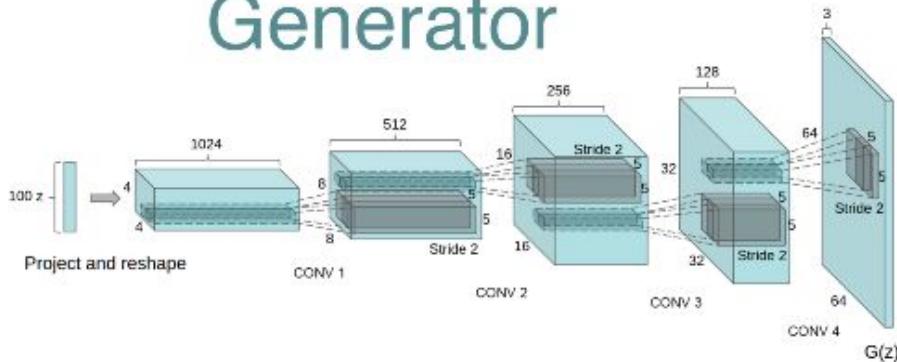
Many other to train G and D :

- f-divergence, Wasserstein loss, feature matching, .... see [19, 20](Jan 2018)

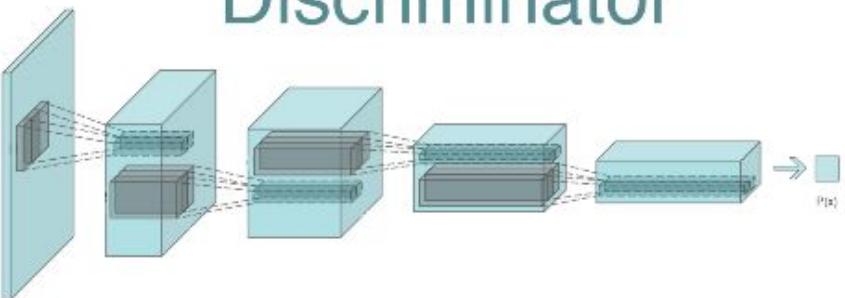
# Which architectures for G and D ?

Ex : Deep Convolutional GAN -DCGAN[21]

## Generator



## Discriminator

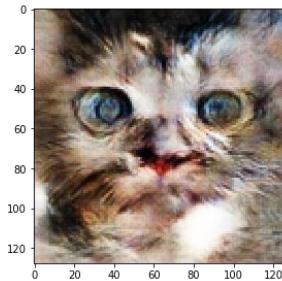
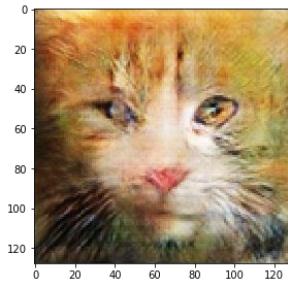
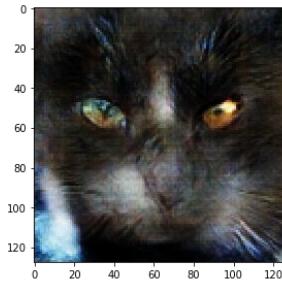
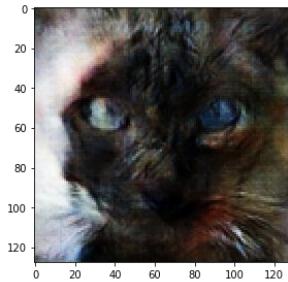


Same improvements as in Style Transfer:

- resized conv > transposed conv
- residual blocks
- several discriminators with random projections [22]

# Some results

GAN, LapGAN, DCGAN, BeGAN, BiGAN,  
DiscoGAN, LSGAN, WGAN, f-GAN,  
Fisher-GAN, AE-GAN, APE-GAN, Gang of  
GANs, InfoGAN, CycleGAN, StackedGAN,  
DualGAN, DeliGAN, .....



-> Meow generator

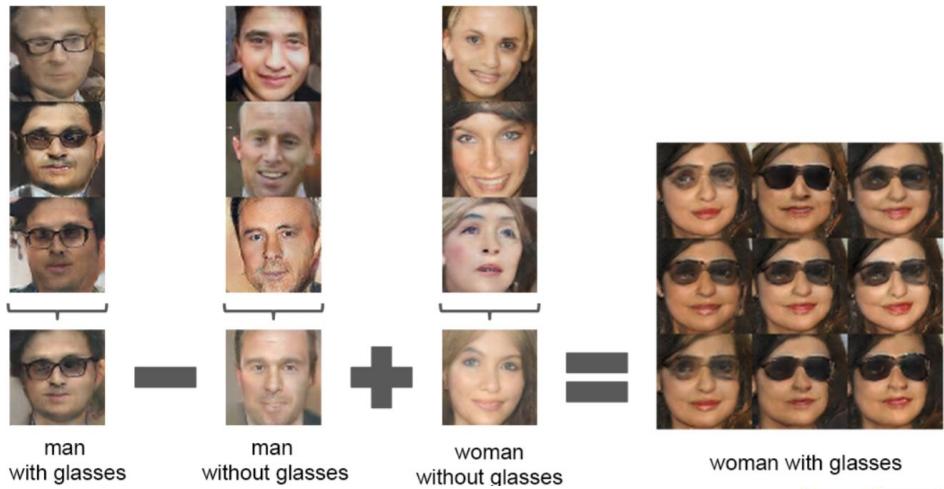
Here, results with a DCGAN, trained with  
'feature matching' loss !



# Latent Space understanding (z)



Arithmetic operation in the latent space :

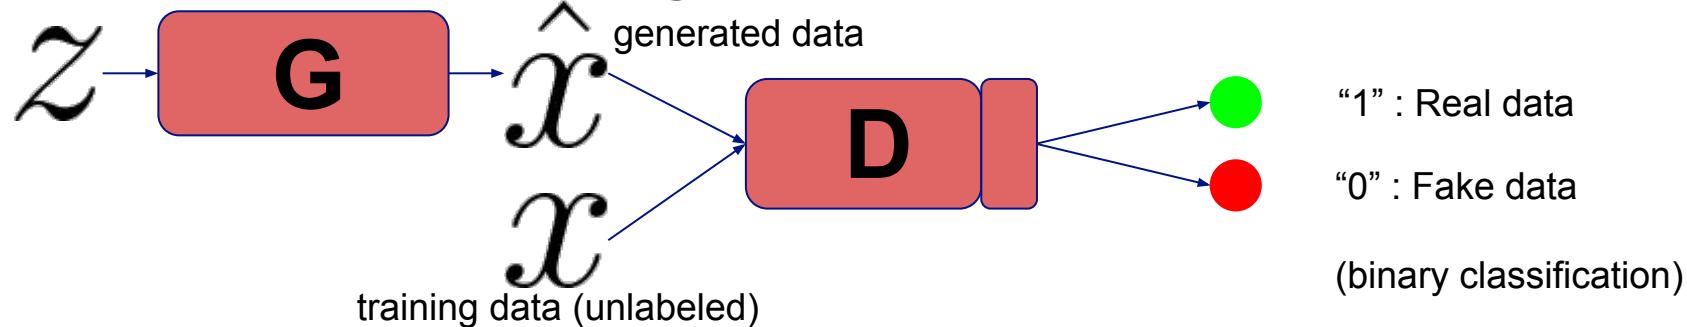


How to get z from a photo :

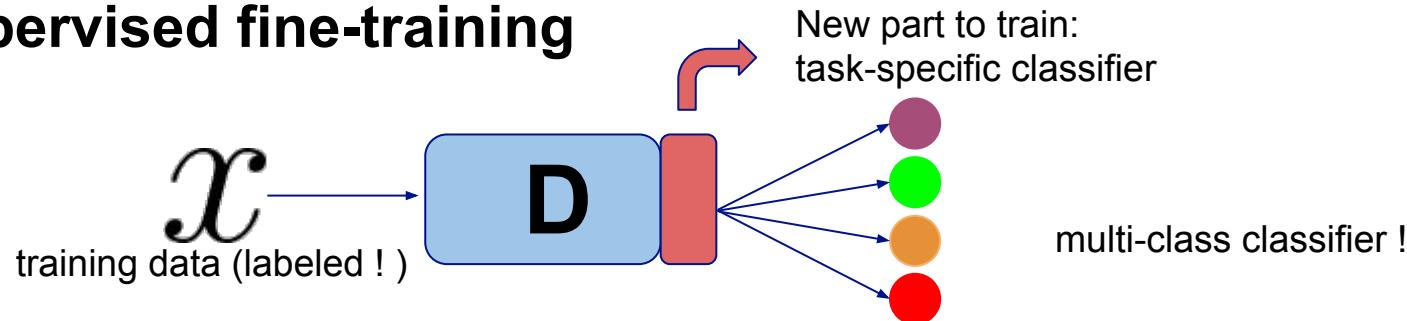
- **recover z by optimization**
- **learn an encoder  $z=E(x)$**  when training D and G
  - BiGAN : GAN + auto-encoder [23]

# GANs for semi-supervised learning

## Unsupervised pre-training



## Supervised fine-training



# Adversarial Domain Adaptation (1/3)

Source domain : SVHN

- with labels



~ 150k samples

**10 classes**, 32x32 pixels

Target domain : MNIST

- without labels



60k + 10k samples

**10 classes**, 28x28 pixels

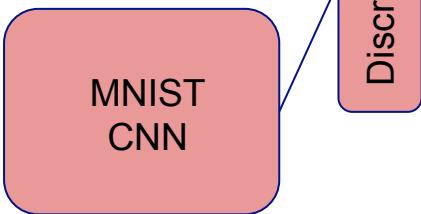
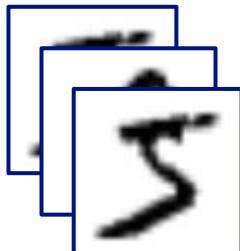
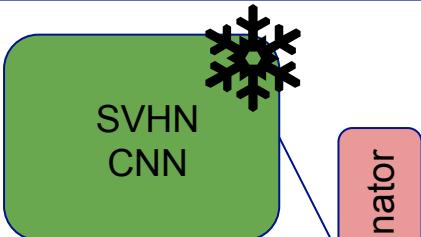
Similar concepts, not the same data source (ex : optical vs SAR images)

# Adversarial Domain Adaptation (2/3)



## Pre-training

- supervised learning
- on source domain
- train 'SVHN CNN' + 'Classifier'

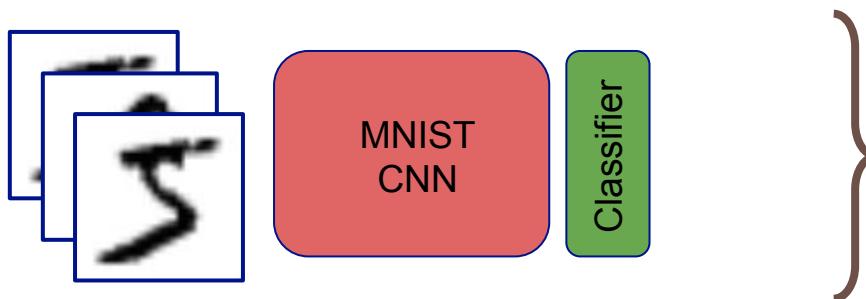


**Task :** binary classification  
- features 'SVHN CNN'  
- or from 'MNIST CNN' ?

## Adversarial Adaptation:

- learn a target encoder CNN (Generator)
- features from 'MNIST CNN' will follow the same distribution as the features from 'SVHN CNN'
- **without labels** from both domains !

# Adversarial Domain Adaptation (3/3)



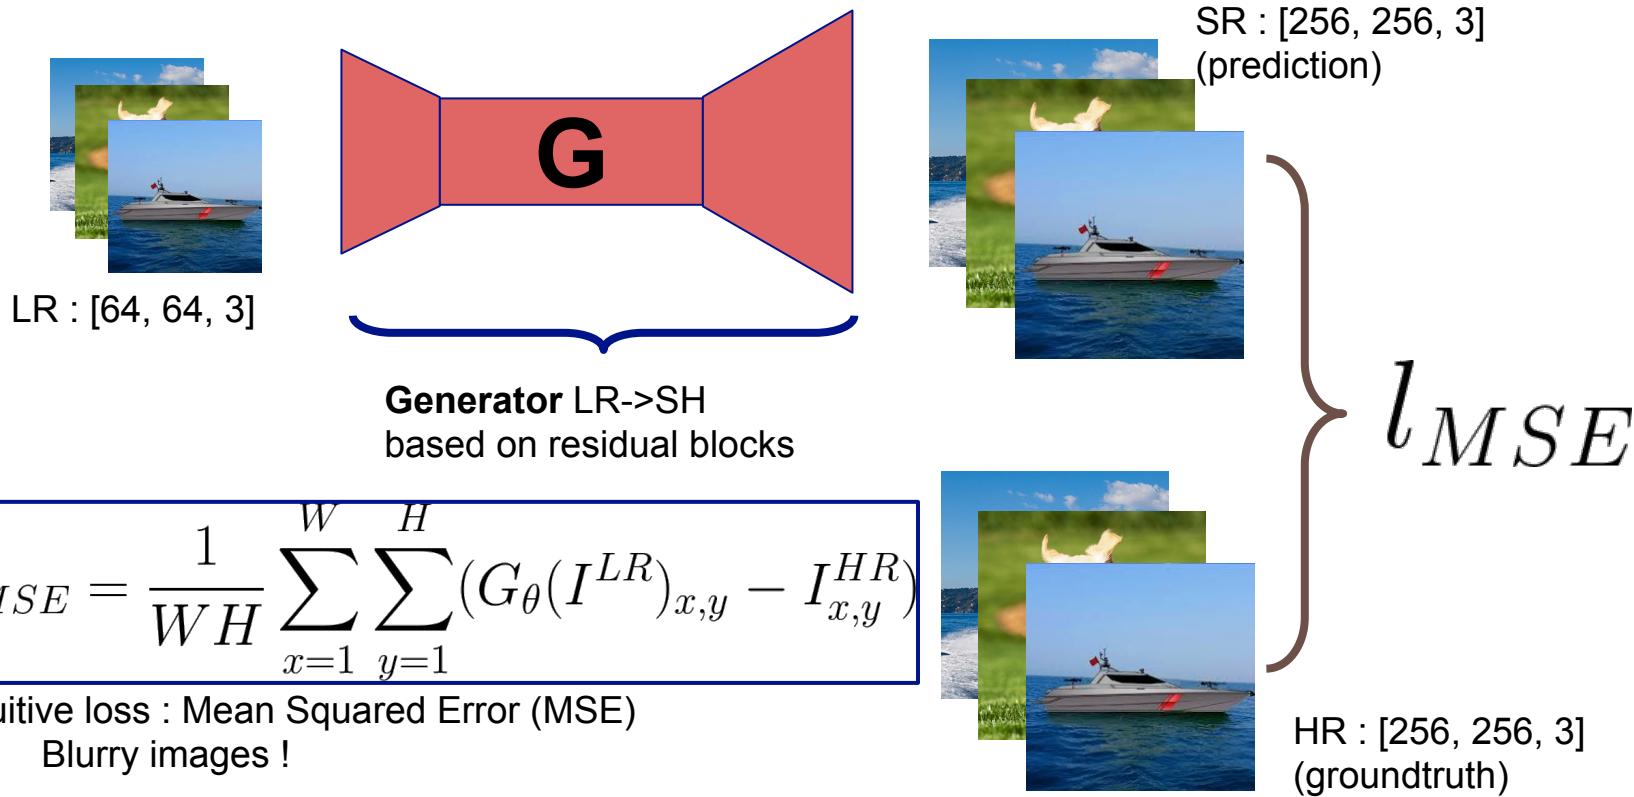
## Testing

- 'Classifier' can understand features from 'MNIST CNN'
- and make classification

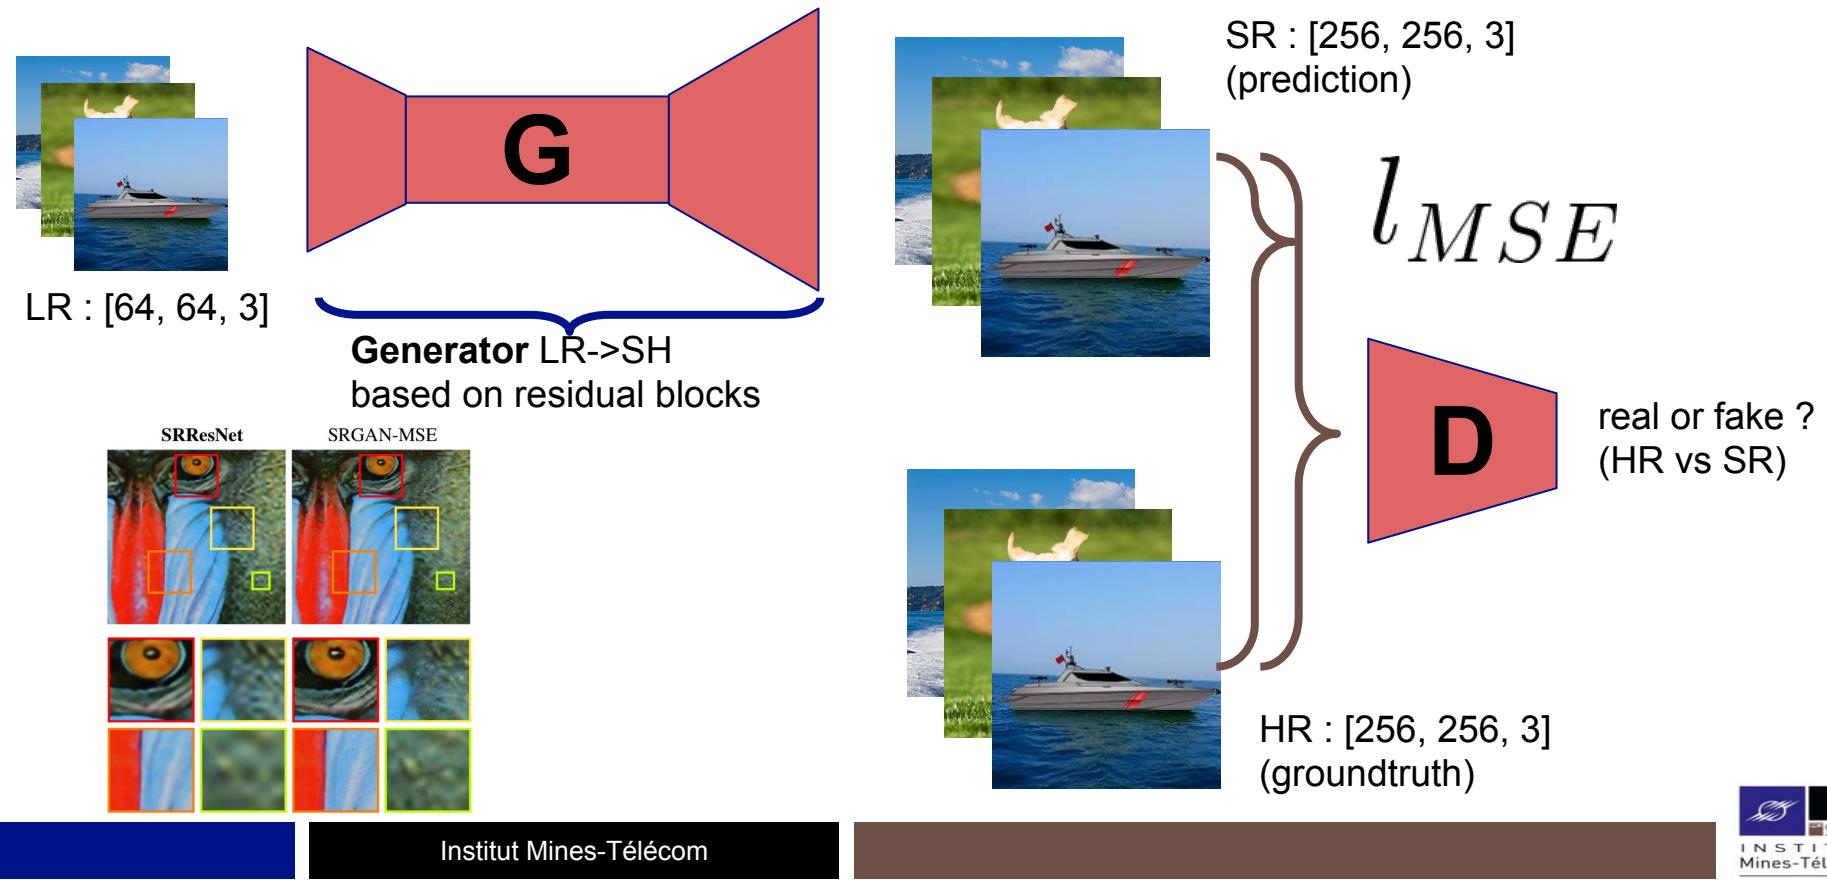
## Results :

Method	MNIST → USPS 173 → 105	USPS → MNIST 105 → 173	SVHN → MNIST 143 51 → 173
Source only	$0.752 \pm 0.016$	$0.571 \pm 0.017$	$0.601 \pm 0.011$
Gradient reversal	$0.771 \pm 0.018$	$0.730 \pm 0.020$	$0.739$ [16]
Domain confusion	$0.791 \pm 0.005$	$0.665 \pm 0.033$	$0.681 \pm 0.003$
CoGAN	$0.912 \pm 0.008$	$0.891 \pm 0.008$	did not converge
ADDA (Ours)	$0.894 \pm 0.002$	$0.901 \pm 0.008$	$0.760 \pm 0.018$

# Enhance Super-Resolution with GANs (1/3)



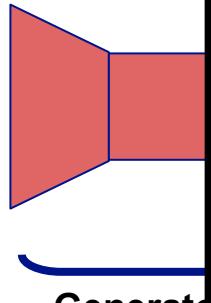
# Enhance Super-Resolution with GANs (2/3)



# Enhance Super-Resolution with GANs (3/3)

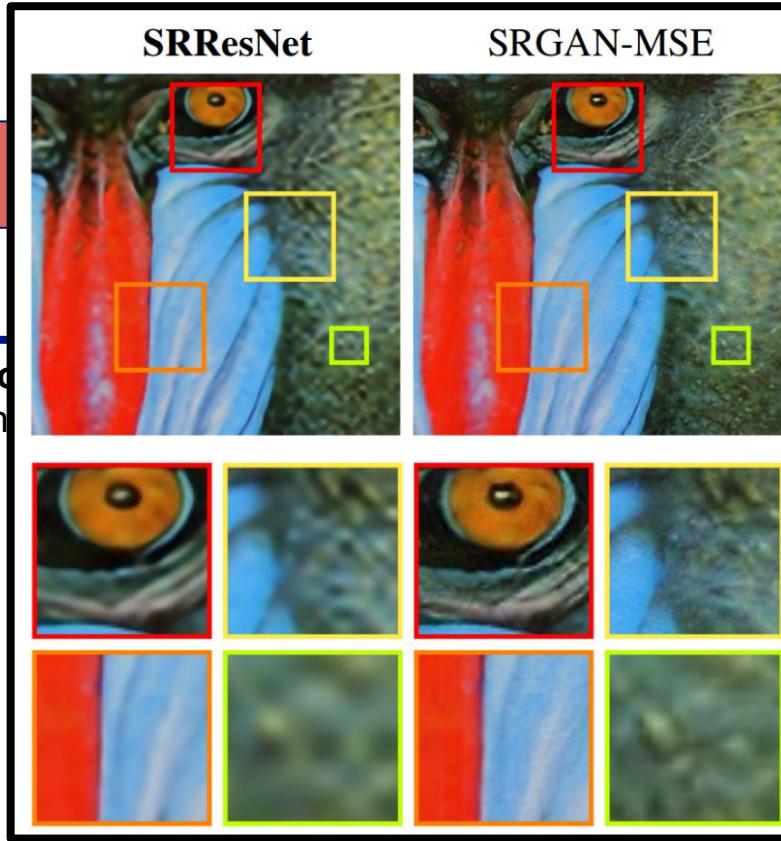


LR : [64, 64, 3]



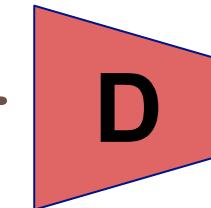
Generator  
based on

[25]



[256, 256, 3]  
diction)

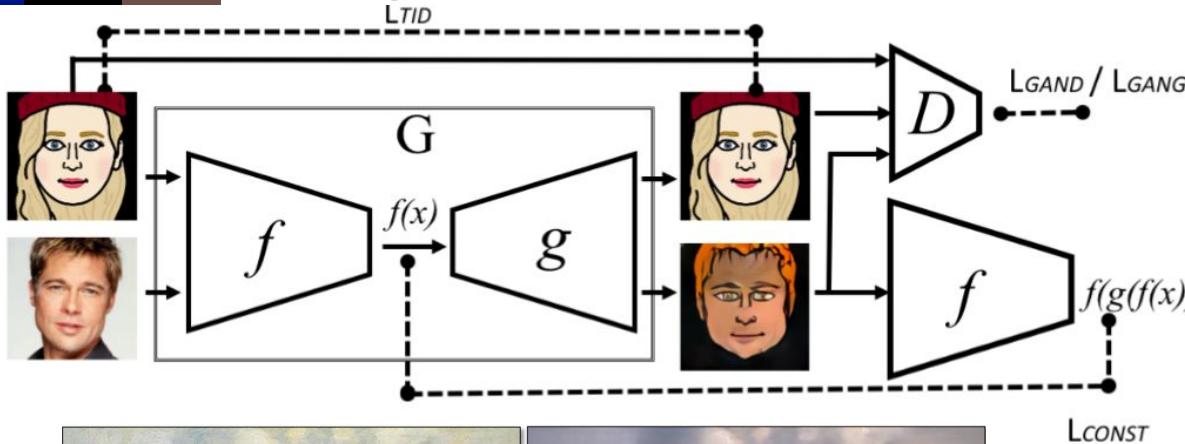
$l_{MSE}$



real or fake ?  
(HR vs SR)

R : [256, 256, 3]  
roundtruth)

# Many applications of GANs ...

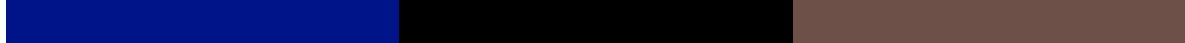


Cross-domain image generation [26] (FAIR)



Reverse style transfer with CycleGAN [27]

paper [28] : "High-Resolution Image Synthesis and Semantic Manipulation with Conditional GANs", Nvidia, Dec 2017  
demo : [https://www.youtube.com/watch?v=3AlpPlzM\\_qs](https://www.youtube.com/watch?v=3AlpPlzM_qs)



# Thanks !

# References (1/2)

- [1] : K. Simonyan, A. Zisserman : “Very Deep Convolutional Networks for Large-Scale Image Recognition”, 2014, [arXiv:1409.1556](https://arxiv.org/abs/1409.1556)
- [2] : C Szegedy et al. : “Going Deeper with Convolutions”, 2014, [arxiv:1409.4842](https://arxiv.org/abs/1409.4842)
- [3] : K He et al : “Deep Residual Learning for Image Recognition”, 2015, [arxiv.org:1512.03385](https://arxiv.org/abs/1512.03385)
- [4] : ImageNet dataset : <http://www.image-net.org/>
- [5] : About Deep Dream visualization technique : [“Inceptionism: Going Deeper into Neural Networks”](#)
- [6] : M. Zeiler, R. Fergus: Visualizing and Understanding Convolutional Networks, 2013, [arXiv:1311.2901](https://arxiv.org/abs/1311.2901)
- [7] : L. Gatys, A. Ecker, M. Bethge : A neural algorithm of artistic style, 2015, [arXiv:1508.06576](https://arxiv.org/abs/1508.06576)
- [8] : M. Ruder, A. Dosovitskiy, T. Brox : Artistic style transfer for video, 2016, [arXiv:1604.08610](https://arxiv.org/abs/1604.08610)
- [9] : Gatys et al : Preserving color in Neural Artistic Style Transfer, 2016, [arXiv:1606.05897](https://arxiv.org/abs/1606.05897)
- [10] : J. Johnson et al : “Perceptual losses for real-time style transfer and super-resolution”, 2016, [arXiv:1603.08155](https://arxiv.org/abs/1603.08155)
- [11] : D. Ulyanov et al : “Instance Normalization: The Missing Ingredient for Fast Stylization”, 2016, [arXiv:1607.08022](https://arxiv.org/abs/1607.08022)
- [12] : MS-COCO dataset : <http://cocodataset.org/#home>
- [13] : V. Dumoulin et al : “A learned representation for artistic style”, 2017, [arXiv:1610.07629](https://arxiv.org/abs/1610.07629)
- [14] : A Aitken et al : “Checkerboard artifact free sub-pixel convolution”, 2017, [arxiv.org:1707.02937](https://arxiv.org/abs/1707.02937)
- [15] : X. Huang and S. Belongie : “Arbitrary Style Transfer in real-time with AdaIN”, 2017, [arXiv:1703.06868](https://arxiv.org/abs/1703.06868)
- [16] : Y Li et al : “Universal Style Transfer via Feature Transforms”, 2017, [arxiv:1705.08086](https://arxiv.org/abs/1705.08086)

# References (2/2)

- [17] : P Isola et al : "Image-to-Image Translation with Conditional Adversarial Networks", 2016, [arxiv:1611.07004](https://arxiv.org/abs/1611.07004)
- [18] : I Goodfellow et al : "Generative Adversarial Networks", 2014, [arxiv:1406.2661](https://arxiv.org/abs/1406.2661)
- [19] : Y Hong et al : "How GANs and its variants work : an overview of GAN", 2017, [arxiv:1711.05914v6](https://arxiv.org/abs/1711.05914v6)
- [20] : S Hitawala : "Comparative Study on GANs", 2018, [arxiv:1801.04271v1](https://arxiv.org/abs/1801.04271v1)
- [21] : A Radford et al : "Unsupervised Representation Learning with Deep Convolutional GANs", 2015, [arxiv:511.06434](https://arxiv.org/abs/1511.06434)
- [22] : B Neyshabur et al : "Stabilizing GAN Training with Multiple Random Projections", 2017, [arxiv:1707.02937](https://arxiv.org/abs/1707.02937)
- [23] : J Donahue et al : "Adversarial Feature Learning", 2016, [arxiv:1605.09782](https://arxiv.org/abs/1605.09782)
- [24] : Eric Tzeng et al : "Adversarial Discriminative Domain Adaptation", 2017, [arxiv:1702.05464](https://arxiv.org/abs/1702.05464)
- [25] : C Ledig et al : "Photo-realistic Single Image Super-Resolution using GANs" 2016, [arxiv:1609.04802](https://arxiv.org/abs/1609.04802)
- [26] : Y Taigman et al : "Unsupervised Cross-Domain Image Generation", 2016, [arxiv:1611.02200](https://arxiv.org/abs/1611.02200)
- [27] : J-Y Zhu et al : "Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks", 2017, [arxiv:1703.10593](https://arxiv.org/abs/1703.10593)
- [28] : T-C Wang et al (NVIDIA) : "High-Resolution Image Synthesis and Semantic Manipulation with Conditional GANs", Dec 2017, [arxiv:1711.11585](https://arxiv.org/abs/1711.11585)